



UNIVERSITY OF PADUA  
UNIVERSITA' DEGLI STUDI DI PADOVA

---

# Zeroth-Order Frank-Wolfe Optimization for Black-Box Adversarial Attacks

Optimization for Data Science, A.Y. 2023/24

---

*Marco Uderzo, 2096998*

- **Goal of the project:** comparing **Zeroth-Order** Frank-Wolfe variants in constraint optimization problems and testing on **black-box adversarial attacks** against MNIST.
- **Implemented Algorithms:**
  - **FZCGS:** Faster Zeroth-Order Conditional Gradient Sliding Method (Gao et al.)
  - **SGFFW:** Stochastic Gradient-Free Frank-Wolfe, with three different gradient approximation schemes:
    - **KWSA:** Kiefer-Wolfowitz Stochastic Approximation
    - **RDSA:** Random Directions Stochastic Approximation
    - **I-RDSA:** Improvised Random Directions Stochastic Approximation
- **Adversarial Attack:** deliberate **manipulation** of the **input** data with the intention of causing a **ML model** to make a mistake or produce **incorrect output**. This is done through **perturbation** of the input in a way that is not easily noticeable to a human observer.

Employs Coordinate-wise Gradient Estimator

$$\hat{\nabla} f(\mathbf{x}) = \sum_{j=1}^d \frac{f(\mathbf{x} + \mu_j \mathbf{e}_j) - f(\mathbf{x} - \mu_j \mathbf{e}_j)}{2\mu_j} \mathbf{e}_j$$

---

**Algorithm 2** Faster Zeroth-Order Conditional Gradient Method (FZCGS)

---

**Input:**  $\mathbf{x}_0, q > 0, \mu > 0, K > 0, \eta > 0, \gamma > 0, n$

```

1: for  $k = 0, \dots, K - 1$  do
2:   if  $\text{mod}(k, q) = 0$  then
3:     Sample  $S_1$  without replacement to compute  $\hat{\mathbf{v}}_k = \hat{\nabla} f_{S_1}(\mathbf{x}_k)$ 
4:   else
5:     Sample  $S_2$  with replacement to compute  $\hat{\mathbf{v}}_k = \frac{1}{|S_2|} \sum_{i \in S_2} [\hat{\nabla} f_i(\mathbf{x}_k) - \hat{\nabla} f_i(\mathbf{x}_{k-1}) + \hat{\mathbf{v}}_{k-1}]$ 
6:   end if
7:    $\mathbf{x}_{k+1} = \text{condg}(\hat{\mathbf{v}}_k, \mathbf{x}_k, \gamma_k, \eta_k)$ 
8: end for

```

**Output:** Randomly choose  $\mathbf{x}_\alpha$  from  $\{\mathbf{x}_k\}$  and return it

---

- $\hat{\nabla} f(\mathbf{x})$  is the estimated gradient of the function  $f$  in  $\mathbf{x}$
- $d$  is the dimensionality of the optimization space
- $\mu_j > 0$  is a smoothing parameter
- $\mathbf{e}_j \in \mathbb{R}^d$  is the basis vector where only the  $j$ -th element is 1 and all others are 0.

---

**Algorithm 3**  $\mathbf{u}^+ = \text{condg}(\mathbf{g}, \mathbf{u}, \gamma, \eta)$  (Qu et al., 2017)

---

```

1:  $\mathbf{u}_1 = \mathbf{u}, t = 1$ 
2:  $\mathbf{v}_t$  be an optimal solution for

```

$$V_{\mathbf{g}, \mathbf{u}, \gamma}(\mathbf{u}_t) = \max_{\mathbf{x} \in \Omega} \langle \mathbf{g} + \frac{1}{\gamma}(\mathbf{u}_t - \mathbf{u}), \mathbf{u}_t - \mathbf{x} \rangle$$

```

3: If  $V_{\mathbf{g}, \mathbf{u}, \gamma}(\mathbf{u}_t) \leq \eta$ , return  $\mathbf{u}^+ = \mathbf{u}_t$ .
4: Set  $\mathbf{u}_{t+1} = (1 - \alpha_t)\mathbf{u}_t + \alpha_t\mathbf{v}_t$  where  $\alpha_t = \min\{1, \frac{\langle \frac{1}{\gamma}(\mathbf{u} - \mathbf{u}_t) - \mathbf{g}, \mathbf{v}_t - \mathbf{u}_t \rangle}{\frac{1}{\gamma}\|\mathbf{v}_t - \mathbf{u}_t\|^2}\}$ .
5: Set  $t \leftarrow t + 1$  and goto step 2.

```

---

- Employs:
  - Coordinate-wise Gradient Estimation
  - Variance Reduction technique
  - Conditional Gradient Sliding

---

**Algorithm 2** Faster Zeroth-Order Conditional Gradient Method (FZCGS)

---

**Input:**  $\mathbf{x}_0, q > 0, \mu > 0, K > 0, \eta > 0, \gamma > 0, n$

```

1: for  $k = 0, \dots, K - 1$  do
2:   if  $\text{mod}(k, q) = 0$  then
3:     Sample  $S_1$  without replacement to compute  $\hat{\mathbf{v}}_k = \hat{\nabla} f_{S_1}(\mathbf{x}_k)$ 
4:   else
5:     Sample  $S_2$  with replacement to compute  $\hat{\mathbf{v}}_k = \frac{1}{|S_2|} \sum_{i \in S_2} [\hat{\nabla} f_i(\mathbf{x}_k) - \hat{\nabla} f_i(\mathbf{x}_{k-1}) + \hat{\mathbf{v}}_{k-1}]$ 
6:   end if
7:    $\mathbf{x}_{k+1} = \text{condg}(\hat{\mathbf{v}}_k, \mathbf{x}_k, \gamma_k, \eta_k)$ 
8: end for
    
```

**Output:** Randomly choose  $\mathbf{x}_\alpha$  from  $\{\mathbf{x}_k\}$  and return it

---



---

**Algorithm 3**  $\mathbf{u}^+ = \text{condg}(\mathbf{g}, \mathbf{u}, \gamma, \eta)$  (Qu et al., 2017)

---

```

1:  $\mathbf{u}_1 = \mathbf{u}, t = 1$ 
2:  $\mathbf{v}_t$  be an optimal solution for
    
```

$$V_{\mathbf{g}, \mathbf{u}, \gamma}(\mathbf{u}_t) = \max_{\mathbf{x} \in \Omega} \langle \mathbf{g} + \frac{1}{\gamma}(\mathbf{u}_t - \mathbf{u}), \mathbf{u}_t - \mathbf{x} \rangle$$

```

3: If  $V_{\mathbf{g}, \mathbf{u}, \gamma}(\mathbf{u}_t) \leq \eta$ , return  $\mathbf{u}^+ = \mathbf{u}_t$ .
4: Set  $\mathbf{u}_{t+1} = (1 - \alpha_t)\mathbf{u}_t + \alpha_t\mathbf{v}_t$  where  $\alpha_t = \min\{1, \frac{\langle \frac{1}{\gamma}(\mathbf{u} - \mathbf{u}_t) - \mathbf{g}, \mathbf{v}_t - \mathbf{u}_t \rangle}{\frac{1}{\gamma}\|\mathbf{v}_t - \mathbf{u}_t\|^2}\}$ .
5: Set  $t \leftarrow t + 1$  and goto step 2.
    
```

---

- It was one of the first approaches to Stochastic Zeroth-Order Frank-Wolfe
- Uses Random Directions Gradient Estimator
- Employs an averaging trick to counter diverging gradient
- Sampled Directions:
  - KWSA: along each coordinate direction (d directions, with d dimensionality of the problem space)
  - RDSA: one random direction
  - I-RDSA:  $m < d$  independently sampled directions

---

## Algorithm 2 Stochastic Gradient Free Frank Wolfe

---

**Require:** Input, Loss Function  $F(x)$ , Convex Set  $\mathcal{C}$ , number of directions  $m$ , sequences  $\gamma_t = \frac{2}{t+8}$ ,

$$(\rho_t, c_t)_{RDSA} = \left( \frac{4}{d^{1/3}(t+8)^{2/3}}, \frac{2}{d^{3/2}(t+8)^{1/3}} \right)$$

$$(\rho_t, c_t)_{I-RDSA} = \left( \frac{4}{(1+\frac{d}{m})^{1/3}(t+8)^{2/3}}, \frac{2\sqrt{m}}{d^{3/2}(t+8)^{1/3}} \right)$$

$$(\rho_t, c_t)_{KWSA} = \left( \frac{4}{(t+8)^{2/3}}, \frac{2}{d^{1/2}(t+8)^{1/3}} \right).$$

**Output:**  $\mathbf{x}_T$  or  $\frac{1}{T} \sum_{t=0}^{T-1} \mathbf{x}_t$ .

1: Initialize  $\mathbf{x}_0 \in \mathcal{C}$

2: **for**  $t = 0, 2, \dots, T-1$  **do**

3:   Compute

    KWSA:

$$\mathbf{g}(\mathbf{x}_t; \mathbf{y}) = \sum_{i=1}^d \frac{F(\mathbf{x}_t + c_t \mathbf{e}_i; \mathbf{y}) - F(\mathbf{x}_t; \mathbf{y})}{c_t} \mathbf{e}_i$$

    RDSA: Sample  $\mathbf{z}_t \sim \mathcal{N}(0, \mathbf{I}_d)$ ,

$$\mathbf{g}(\mathbf{x}_t; \mathbf{y}, \mathbf{z}_t) = \frac{F(\mathbf{x}_t + c_t \mathbf{z}_t; \mathbf{y}) - F(\mathbf{x}_t; \mathbf{y})}{c_t} \mathbf{z}_t$$

    I-RDSA: Sample  $\{\mathbf{z}_{i,t}\}_{i=1}^m \sim \mathcal{N}(0, \mathbf{I}_d)$ ,

$$\mathbf{g}(\mathbf{x}_t; \mathbf{y}, \mathbf{z}_t) = \frac{1}{m} \sum_{i=1}^m \frac{F(\mathbf{x}_t + c_t \mathbf{z}_{i,t}; \mathbf{y}) - F(\mathbf{x}_t; \mathbf{y})}{c_t} \mathbf{z}_{i,t}$$

4:   Compute  $\mathbf{d}_t = (1 - \rho_t) \mathbf{d}_{t-1} + \rho_t \mathbf{g}(\mathbf{x}_t, \mathbf{y}_t)$

5:   Compute  $\mathbf{v}_t = \operatorname{argmin}_{\mathbf{s} \in \mathcal{C}} \langle \mathbf{s}, \mathbf{d}_t \rangle$ ,

6:   Compute  $\mathbf{x}_{t+1} = (1 - \gamma_t) \mathbf{x}_t + \gamma_t \mathbf{v}_t$ .

7: **end for**

---

- The experiment conducted on these **Frank-Wolfe** variants is a **Black-Box Adversarial Attack on the MNIST Dataset**.
  - This means that we have to find the **smallest perturbation** of the input 28x28 greyscale image of a handwritten digit, constrained by  $\|\delta\|_{\infty} \leq s$ , able to make the Deep Neural Network take the incorrect prediction.
- We follow the **setup** from `nn-carlini` pre-trained DNN for the MNIST dataset.
- **Feasible Set** used:
  - As per section 4.3 of Gao et al., the feasible set used is  $\|\delta\|_{\infty} \leq s$
  - $s$  has different **values** based on the algorithm:
    - **FZCGS**:  $s = 4$
    - **SGFFW** with **RDSA**:  $s = 8$
    - **SGFFW** with **I-RDSA** and **KWSA**:  $s = 4$

- The implementative changes from the Theoretical Papers concern the **FZCGS algorithm** (Gao et al.). In particular, the **Conditional Gradient Procedure**.
- $\text{condg}()$  from Gao et al. was **changed** with the  $\text{condg}()$  version from the paper Lobanov et al.
- $\text{condg}()$  also presents an **additional stopping condition** based on the value of the **alpha** parameter.
  - This additional condition **speeds up** the already extremely **long running times** of FZCGS, whilst producing **negligible deviations** from the otherwise computed result.

---

**Algorithm 3**  $u^+ = \text{condg}(g, u, \gamma, \eta)$  (Qu et al., 2017)

---

- 1:  $u_1 = u, t = 1$
- 2:  $v_t$  be an optimal solution for

$$V_{g,u,\gamma}(u_t) = \max_{x \in \Omega} \langle g + \frac{1}{\gamma}(u_t - u), u_t - x \rangle$$

- 3: If  $V_{g,u,\gamma}(u_t) \leq \eta$ , return  $u^+ = u_t$ .
  - 4: Set  $u_{t+1} = (1 - \alpha_t)u_t + \alpha_t v_t$  where  $\alpha_t = \min\{1, \frac{\langle \frac{1}{\gamma}(u - u_t) - g, v_t - u_t \rangle}{\frac{1}{\gamma}\|v_t - u_t\|^2}\}$ .
  - 5: Set  $t \leftarrow t + 1$  and goto step 2.
- 

Conditional Gradient procedure (Gao et al.)

---

**Algorithm 2** Conditional Gradient procedure

---


- 1: for  $t = 0, \dots, T$  do
  - 2:    $v_t \leftarrow \text{argmin}_{v \in Q} \langle g_t, v \rangle$
  - 3:   if  $\langle g_t, u_t - v_t \rangle \leq \beta$  then
  - 4:     return  $u_t$
  - 5:   end if
  - 6:    $\alpha_t \leftarrow \min\left\{\frac{\langle g_t, u_t - v_t \rangle}{\eta\|u_t - v_t\|^2}, 1\right\}$
  - 7:    $u_{t+1} \leftarrow u_t + \alpha_t(v_t - u_t)$
  - 8:    $g_{t+1} \leftarrow g_0 + \eta(u_{t+1} - u_0)$
  - 9: end for
- 

Conditional Gradient procedure (Lobanov et al.)

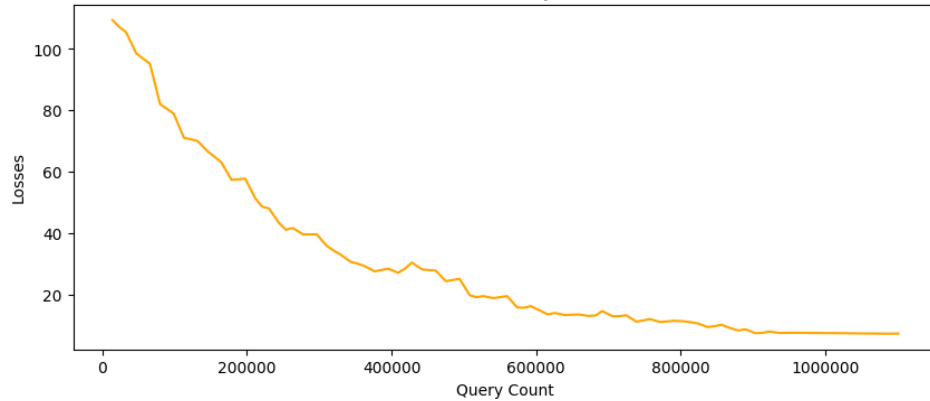
# FZCGS Results




Losses vs. Iteration Index



Losses vs. Query Count



Loss vs. Time








# SGFFW Results (Compared)



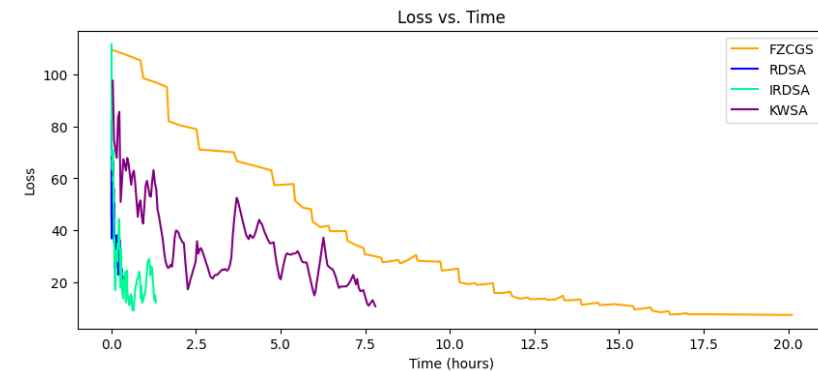


- **SGFFW** has different behaviours depending on which gradient approximation scheme is used.
- **RDSA**: fastest in time and with limited query count, while requiring a lot of iterations.
- **I-RDSA**: requires extremely less iterations than RDSA, but requires more queries.
- **KWSA**: slowest in time, requires the least amount of iterations, at the expense of a very large query count.
- **I-RDSA** shows a remarkable **compromise** between **iterations**, **query count** and **running time** to converge.




# Results: All Algorithms Compared



| Algorithm    | Metric             | Value   |
|--------------|--------------------|---------|
| FZCGS        | Best Loss          | 7.283   |
|              | Last Loss          | 7.312   |
|              | Iteration Count    | 100     |
|              | Query Count        | 1100736 |
|              | Running Time (Hrs) | 20.1    |
| SGFFW-RDSA   | Best Loss          | 16.9    |
|              | Last Loss          | 18.922  |
|              | Iteration Count    | 6000    |
|              | Query Count        | 12000   |
|              | Running Time (Hrs) | 0.43    |
| SGFFW-I-RDSA | Best Loss          | 8.96    |
|              | Last Loss          | 12.156  |
|              | Iteration Count    | 500     |
|              | Query Count        | 55000   |
|              | Running Time (Hrs) | 1.3     |
| SGFFW-KWSA   | Best Loss          | 10.647  |
|              | Last Loss          | 10.647  |
|              | Iteration Count    | 200     |
|              | Query Count        | 313600  |
|              | Running Time (Hrs) | 7.7     |





- There is **no overall best algorithm**, since each of them has **strengths** and **weaknesses**.
- **FZCGS** has the **best** performance for **iterations** after SGFFW with **KWSA**, but is the **worst** for query count and running times.
- **SGFFW with I-RDSA** shows a remarkable **compromise** between all metrics:
  - Requires an acceptable amount of **iterations**
  - The **query** count is reasonably small
  - Running **times** are convenient
  - Reaches one of the **lowest losses**